Voice Input Adoption (ChatGPT India Mobile)

Turning voice from something Bharat users have tried once into something they keep using.

Team: ChatGPT Mobile Team (India)

Contributors: Mayank Sharma (PM Fellow, Nextleap) — PM/author.

Status: In Review

Launching on: To Be Decided — gated on baseline instrumentation and a 4-arm experiment. Target decision point: W10.

Resources: Milestone1 - market sizing, competitive teardown, North Star definition · Milestone2 - segment scorecard, user research, Problem Framing Canvas, User flow (FigJam), Annotated wireframes (Figma), Survey instrument and raw responses: Survey Response Excel Sheet

Problem Definition

Voice input on ChatGPT mobile India is under-used despite India being a mobile-first, voice-native market, which is second-largest market at 100M+ WAU, where 98% of internet users consume Indic-language content and WhatsApp alone carries roughly 7B voice messages a day. The users most affected are Vernacular "Bharat" users: Tier-2/3, regional-language-first Indians for whom typing in Devanagari or transliterating into Hinglish is the single largest tax on getting an answer. Our research shows the failure is not awareness, 87% of surveyed users had already tried voice at least once, rather, it is that the first attempt never converts into a habit, so users try, hit friction, and quietly revert to typing. If we solve it, those users get answers at speaking speed with no script barrier, and the business unlocks the largest reachable base in the market (500M+ addressable, no entrenched AI rival) plus retention and monetisation headroom in a segment we currently under-serve.

Question	Answer
What is the problem?	Voice input adoption is low. Trial is high, 87% have tried, but repeat is not. Users try once, hit friction, and revert to typing. The gap is habit formation, not awareness.
Who is facing it?	Vernacular "Bharat" users, Tier-2/3, regional-language-first Indians. ~500M+ addressable. Selected over five alternative segments in Milestone 2.
What business value is unlocked?	The largest reachable base in India with no entrenched AI rival. Retention and monetisation headroom in an under-served segment. A defensive position against Gemini Live on the voice axis.
How do target users benefit?	They stop paying the typing tax. A Hindi-first user who types Devanagari slowly or transliterates awkwardly into Hinglish, asks at speaking speed instead. Voice is also script-agnostic, which removes the literacy barrier entirely for users who speak far better than they type.
Why is it urgent now?	Gemini Live is positioning on exactly this axis. India is simultaneously at peak voice-readiness: 98% Indic-content consumption, ~7B WhatsApp voice messages/day, 140M existing voice-command users, world's cheapest mobile data. The habit is forming right now with someone.

Goals

High-level goals in priority order, then the measurable metrics that test each.

1. **G1 - Convert first-time voice users into repeat voice users.** The habit goal. This is the primary objective: our own data says trials already happen and does not stick.

2. **G2 - Make the mic discoverable enough that a first try happens at all.** The breadth goal. Secondly, because the trial may be less broken than assumed, see the sampling caveat.
3. **G3 - Make a wrong word cheap to fix, so one error does not end the relationship.** The recovery goal. Supports G1: the most common reason a first try fails to become a second try is an uncorrectable mistake.

How these ladder to the North Star. Milestone 1 defined the North Star as **Total Voice Inputs / week**, decomposed into **Active Voice Users × Voice Inputs per Active User**. G2 drives the first term (breadth). G1 and G3 drive the second (depth). Direction A is deliberately chosen to move both.

Functional metrics

Metric	Definition	Why it matters
voice repeat rate - PRIMARY	% of users who send a voice input and send another within 7 days	Tests G1. Trial is already high and repeat is not, so this isolates habit formation, the only thing this feature actually attempts to change. Activation metrics would flatter us.
Voice share of inputs	Voice inputs ÷ all inputs, per user	Confirms a durable mix shift rather than a novelty spike. A streak can inflate repeat rate for two weeks; share is harder to fake.
Mic tap rate	Mic taps ÷ weekly active users	Tests G2 in isolation. If this stays flat, the promoted mic and nudge simply are not being seen.
Voice start → send completion rate	Voice sends ÷ capture starts	Catches abandonment mid-dictation - the signature of an ASR or dictation-UX failure rather than a discoverability one.
Transcript edit rate	% of voice sends where at least one word was corrected	Tests G3 and doubles as a diagnostic. Rising = recognition degrading. Near zero = the correction chips are unnecessary and can be cut.

Non-functional metrics (guardrails that must not regress)

Guardrail	Why it matters	Trip condition
Total inputs per user	Cannibalisation check. The most likely way this feature fails quietly.	Voice inputs up while total inputs flat
Nudge dismiss rate	Fatigue signal. We are shipping an interruption.	Rising over time
Session abandonment after nudge	Are we interrupting people out of the app?	Any increase vs control
Typed-input completion time	Did the promoted mic clutter the composer for people who prefer typing?	Any regression vs control
Mic permission denial rate	Are we asking for the mic too early in onboarding?	Increase vs control
app retention	Blunt safety net for anything the above misses.	Any decrease vs control

Non-Goals

Explicitly out of scope. Anything here is either owned elsewhere, deferred, or deliberately rejected.

- **Improving speech-recognition accuracy.** ASR word-error-rate is an Engineering / Applied-AI workstream. We track WER by language as a dependency and report it, but we never claim it as a deliverable of this project. This scope line is load-bearing, see Trade-offs.
- **Advanced Voice Mode paywall or pricing changes.** Direction D territory; 74% have never seen the free-tier limit message, so the addressable population is small.
- **Voice output / text-to-speech.** This project is about input only.
- **Desktop and web voice input.** Mobile India only.
- **A discreet or "whisper" input mode.** Direction C considered and descoped; requires mic-sensitivity work we do not own, and the barrier is partly environmental.
- **Native-script keyboard or transliteration improvements.** A different solution to the same typing tax; not this bet.
- **Offline or low-connectivity voice.** Hypothesis H10 territory; deferred.
- **Segments outside Vernacular Bharat.** Students, Urban Professionals and Developers are deprioritised. Gig & Field Workers are Phase 2. Older / Low-Literacy Adults are Phase 3.

Validation of the problem

User research (condensed survey, n = 30+) ChatGPT mobile users in India

Each figure traces to exactly one survey question. No number is reused for two meanings; no categories are silently combined.

What we asked	What we found	What it implies
Have you tried voice input?	87% yes	Awareness and trial are not the primary gap. This is the finding that reframed the whole project.
Rate recognition accuracy (1-5)	3.45 average	Mediocre, not catastrophic. Bad enough to annoy, good enough that the product is not the blocker for everyone.
What problems did you face?	48% wrong words (most-cited) · 23% Hinglish mixing · 19% accent/dialect · 19% background noise · 23% no problems at all	Errors are common but not universal. The 23% "no problems" is genuine counter-evidence and is kept visible deliberately.
How comfortable are you speaking aloud near others?	13% very comfortable · 39% with family only · 29% only when alone · 19% would never	48% are socially constrained. A nudge-based solution has a real downside for roughly half the segment.
Have you hit the free-tier voice limit?	74% never saw the message · 13% stopped using voice after it · 6% switched to typing · 6% considered upgrading	The paywall is a narrow problem affecting a small slice. Deprioritised on that basis.
Open-ended: what stops you using voice? (25 of 30+ answered)	8 of 25 prefer typing / habit · 5 of 25 accuracy · 2 of 25 social discomfort · 1 of 25 paywall cost	Habit is the largest single theme and larger than accuracy. This is the direct evidence for Direction A.

The decisive pattern: trial is high, repeat is low, and the most common self-reported reason is habit rather than capability. That is what makes this a product problem and not purely a model problem and it is why the primary metric is a repeat metric, not an activation metric.

Competitive landscape

The Milestone 1 teardown found ChatGPT trailing Gemini Live, WhatsApp, Perplexity and Wispr Flow on four axes. Two of those four are product-owned, which is exactly the argument for this project existing.

Axis	Status vs competitors	Owner
Voice entry-point prominence	Behind: our mic is a small icon; rivals surface voice as a primary action	Product in scope (Screens 1–2)
Transcript editability	Behind: rivals let you fix a word; we make you re-record	Product in scope (Screen 4)
Code-switching (Hinglish) tolerance	Behind	Applied AI dependency, not in scope
Latency	Behind	Engineering dependency, not in scope

Market context (Milestone 1)

- 806–886M internet users; 98% consume Indic-language content (IAMAI-Kantar, 2024/2025)
- 140M existing voice-command users, 57% of them rural; the behaviour already exists, just not in our app
- WhatsApp carries roughly 7B voice messages/day; voice is already a social norm, not a novelty
- ASR word-error-rate benchmarks: Hindi ~16%, low-resource languages 35%+; this is our ceiling, and we do not own it
- 62% of Indian users report AI-privacy concerns (BCG); relevant to mic-permission timing

Understanding the target audience

Segment in focus and size

Vernacular "Bharat" users: Tier-2/3, regional-language-first Indians. **~500M+ addressable.** Selected in Milestone 2 over five alternatives using a scorecard of Segment Size, Voice Affinity, Product-Fixable barriers, and Business Impact.

Segment	Decision and reasoning
Vernacular Bharat	**SELECTED**: largest reachable base with no strong AI rival; highest natural voice affinity (already voice-first on WhatsApp and Google); script-agnostic advantage; clear product-owned levers.
Students & Exam Aspirants	**Deprioritised**: largest ChatGPT cohort globally, but the barrier is environmental (silent study spaces), not product-fixable.
Urban Professionals	**Deprioritised**: low voice affinity; prefer typing for precision. Task mismatch.
Gig & Field Workers	**Phase 2**: very high voice affinity and genuine hands-free need, but a smaller base.
Developers & Builders	**Deprioritised**: voice is fundamentally the wrong tool for code.
Older / Low-Literacy Adults	**Phase 3**: highest dependency on voice, but lowest current ChatGPT reach and hardest to onboard.

Footnote carried from Milestone 2: speech-recognition (ASR) quality is an Engineering / Applied-AI-owned workstream, tracked as a dependency and not claimed as a product deliverable. "Product-Fixable" counts only onboarding, discoverability and pricing/UX levers this team owns directly.

Key personas

Persona 1: Ramesh, 29, Bhopal. Small trader, part-time farming.

- **Language:** thinks and speaks Hindi; reads English slowly; types Devanagari painfully, so transliterates into Hinglish.
- **Behaviour:** already voice-first; sends WhatsApp voice notes daily, uses Google voice search without thinking about it.
- **Device and plan:** mid-tier Android, free ChatGPT tier, cheap abundant data.
- **Context and queries:** practical and long; fertiliser quantities, crop timing, government forms, GST questions. His queries are exactly the ones where typing hurts most.
- **Why he matters:** he is the clearest case for the contextual nudge; long queries, high typing friction, no social barrier when working alone.

Persona 2: Kavita, 34, Sagar. Runs a small tailoring business from home.

- **Language:** Hindi only. Low English comfort. Speaks far better than she types; the literacy gap is the whole problem.
- **Behaviour:** voice notes constantly on WhatsApp, but shares her phone with family and lives in a full house.
- **Device and plan:** entry-level Android, shared device, free tier.
- **Context and queries:** pricing, measurements, small-business and health questions; some of them private.
- **Why she matters:** she is the counter-case. She has the highest need for voice and the highest social cost of using it. She represents the 29% who use voice only when alone and the 19% who never would. If our nudge annoys Kavita, the guardrails should catch it, she is the reason the frequency governor exists.

User journey: today (as-is)

Opens the app → sees a text field → types slowly in Hinglish, backspacing repeatedly → gets a useful answer → either never registers the mic at all, or tried it once, got a wrong word, found re-recording worse than typing, and quietly went back. **Nothing in the product ever invites a second attempt.** That dead end is the problem in one sentence.

User journey: proposed (to-be)

Opens the app → during language selection, sees a one-time voice teaser in her own language (Screen 1) → in chat, the mic sits in the composer's primary slot with a label she can read (Screen 2) → mid-way through a long typed query, a coach mark offers voice at the exact moment typing hurts (Screen 3) → speaks; the transcript appears live and one wrong word is fixed with a single tap rather than a re-record (Screen 4) → sends; a light confirmation and streak nudge give her a reason to come back tomorrow. **Full flow:** [User flow (FigJam)]().

Unmet needs

Goals

- Get a useful answer without paying the typing tax.
- Ask in the language they actually think in, without fighting a script.
- Trust that a mistake is recoverable before committing to a new input method.

Pain points

- The first voice attempt produced a wrong word, and re-recording the whole query felt worse than typing it. (48% wrong-word errors)
- Speaking aloud is socially costly. (48% constrained to family-only, alone-only, or never)
- The mic is not visible enough to become a default. (Competitive teardown: entry-point prominence)
- Nothing rewards the second try, one bad attempt becomes a permanent verdict. (8 of 25 now simply prefer typing)

Solution

Direction chosen: habit-reframing onboarding and mic discoverability, with a minimal correction layer folded in. It attacks the largest barrier we discovered in our own data i.e **habit** - which is also the barrier most within a PM's power to move: habit is shaped by nudges, defaults and reinforcement, not by fixing a model. It carries no hard engineering dependency, so it ships and A/B-tests independently of the ASR roadmap.

Directions considered

Direction	Impact	Effort	Confidence	Evidence anchor and decision
A. Habit-reframing onboarding & mic discoverability	High	Low	High	Habit is the largest qualitative theme (8 of 25). Fully product-owned. **SELECTED.**
B. Correction-first voice UX	High	Medium	Med-High	48% reported wrong-word errors, the top technical complaint. **Partially folded in** (tap-to-fix chips, Screen 4); full scope is Phase 2 because it depends on ASR confidence-score access.
C. Discreet / "whisper" mode	Medium	High	Medium	48% socially constrained. **Descoped** needs mic-sensitivity work owned by Applied AI; the barrier is partly environmental.
D. Free-tier paywall friction	Low-Med	Low	Medium	74% never saw the limit message; 13% stopped after it. **Descoped** small addressable population; handled as a separate quick win.

How the solution draws on prior learnings

- **From user research (habit):** the nudge fires on real typing friction rather than at random. We interrupt at the moment typing hurts, which is the only moment a habit is actually re-negotiable.
- **From user research (socially constrained):** the nudge is cheap to dismiss and governed by a frequency cap. We deliberately do not push voice at people who have told us it is socially costly.
- **From user research (wrong-word errors):** tap-to-fix chips instead of re-record. This is the cheapest slice of Direction B, and it exists because an uncorrectable error is the most likely reason a first try never becomes a second.
- **From competitive research:** entry-point prominence and transcript editability are the two product-owned gaps of the four identified. Screens 2 and 4 map onto them one-to-one.
- **From Milestone 1 (North Star decomposition):** features are deliberately split across both terms, discoverability drives Active Voice Users, the reinforcement loop drives Inputs per Active User.
- **From Milestone 2 (segment scorecard):** every feature here sits inside the "Product-Fixable" column — onboarding, discoverability, UX. Nothing claims ASR.

User flows, wireframes and mockups

- **User flow** (FigJam): [User flow (FigJam)]
- **Annotated wireframes**, 4 screens with 12 annotations (Figma): [Annotated wireframes (Figma)]

Screen	State	Intervention and rationale
1	Onboarding, first run	Language-first voice teaser. Introduces voice at the exact moment the user declares she thinks in Hindi; peak relevance. Skippable; no forced adoption.
2	Chat home, returning user	Mic was promoted to the composer's primary action slot with a label in the user's language. Past voice messages tagged "sent by voice" to normalise the behaviour.
3	Mid-typing, long query	Contextual coach mark framed on speed ("3x tez"), not features. Cheap dismisses that reduce future frequency.
4	Dictation active	Live editable transcript; low-confidence word underlined with one-tap correction chips; first-use confirmation and streak nudge.

Key features (the user benefits to be built)

1. **Language-first onboarding voice teaser** (Screen 1). *Benefit:* the user learns that a voice exists in her language at the one moment she is thinking about language. *Serves:* G2.
2. **Promoted mic in the composer** (Screen 2). *Benefit:* the mic stops being a hidden icon and becomes an obvious choice, labelled in a language she can read. *Serves:* G2.
3. **Contextual nudge on typing friction** (Screen 3). *Benefit:* the suggestion arrives when typing is actively painful, not as generic promotion. *Serves:* G1.
4. **Live editable transcript with tap-to-fix** (Screen 4). *Benefit:* one wrong word costs one tap instead of re-speaking the whole query. *Serves:* G3.
5. **First-use reinforcement** (Screen 4). *Benefit:* a reason to come back tomorrow. *Serves:* G1 and this is the loop that moves Inputs per Active User.

Key logic (algorithm, schema and data changes)

- **Friction detection (new client-side heuristic).** Triggers on query length above a threshold, backspace density, and grip/orientation signal. Purely local; no new data types; no server round-trip.
- **Nudge frequency governor (new per-user state).** Dismissal counter persisted per user; cap tightens after two dismissals. Small schema addition: one counter, one timestamp.
- **Correction chips (no model change).** Consume ASR confidence scores and n-best alternates **already returned** by the recognition service. This is a dependency on Applied AI's existing output, not a request to improve recognition. If those fields are not currently exposed, that is a blocking prerequisite (see Open Questions).
- **Streak state (new per-user counter).** One counter, one last-used timestamp.
- **Segment targeting (derived attribute).** Experiment bucketing on city tier × app language. This is a proxy for the Bharat segment, not a clean identifier and stated as such rather than implied.
- **Privacy posture: unchanged.** No new PII collected; voice audio handling and retention are untouched by this project.

Instrumentation (events to be logged)

- voice_prompt_shown { surface: onboarding | nudge, trigger, language, session_n }
- voice_prompt_dismissed { surface, lifetime_dismiss_count }
- mic_tapped { entry: composer | nudge | onboarding }
- voice_capture_started · voice_capture_cancelled { duration_ms }
- transcript_shown { asr_confidence, language_detected, word_count }
- transcript_edited { method: chip | manual, word_index }
- voice_input_sent { duration_ms, word_count, was_edited, language }
- voice_streak_shown { streak_day }

Cut against: city tier, app language, device tier, connectivity.

Launch Readiness

Key milestones

Week	Milestone	Exit criteria
W0-W2	Instrumentation ships; baseline logged	A real voice repeat baseline exists for the first time. Sampling question (87% trial) resolved against logged behaviour.
W2	Design complete	Composer, coach-mark and transcript specs signed off against the ChatGPT design system. Targets and MDE set from baseline; not before.
W3-W5	Development complete	All four experiment arms built behind feature flags.
W6	QA and dogfooding	Tested on mid-tier and entry-level Android, on-device, with India-based Hindi and regional-language speakers. Hinglish code-switching cases explicitly in the test plan.
W7-W9	Experiment live	Minimum 2 weeks so that two D7 windows are observed. No early peeking.
W10	Readout and decision	Ship / iterate / kill against the pre-committed criteria below.

Launch checklist and internal stakeholders

Stakeholder	The question they will ask	Action required
Applied AI / Speech	"Are you asking us to improve ASR?"	No. Confirm confidence scores and n-best alternates are exposed on the existing API. Agree a WER-by-language reporting cadence so we can read our own ceiling.
Design	"Does a coach mark fit our design language?"	Composer and coach-mark specs; review the mic promotion against the existing composer hierarchy.
Localisation	"Is the copy Hinglish in Latin script, or native Devanagari?"	Copy review and a script decision. Currently unresolved. (see Open Questions).
Legal / Privacy	"When do we ask for the mic? What do we say about voice data?"	Permission-prompt timing and privacy copy. 62% of Indian users report AI-privacy concerns, so this is not a formality.

Stakeholder	The question they will ask	Action required
Support / Ops	"What will users contact us about?"	FAQ for the nudge and the streak; expected contact drivers; a documented way to turn the nudge off.
Data / Analytics	"Is the event schema sound? Is the segment proxy defensible?"	Schema review; sign-off on Tier-2/3 × app-language as the Bharat proxy, with its limitations written down.
Marketing / Comms	"Is there an Indian voice story here?"	Decide launch vs silent rollout. Recommend silence until the experiment reads out.

Experimentation plan

Four arms, user-level randomisation. Population: Tier-2/3 × regional app language.

Arm	Treatment	What it isolates
1. Control	Current experience	Baseline
2. Discoverability	Onboarding teaser + promoted mic (Screens 1–2)	G2. Does anyone try a voice who is not going to?
3. Habit loop	Contextual nudge + reinforcement + tap-to-fix (Screens 3–4)	G1 and G3. Do people who already tried voice come back?
4. Full bundle	All four screens	Interaction effects between the two

Primary metric: voice repeat rate. **Unit:** user. **Duration:** ≥ 2 weeks. **Why four arms rather than a simple A/B:** a bundled test tells us that it worked, not which part did. Given 87% trial, the working hypothesis is that Arm 3 carries most of the lift and Arm 2 barely moves. Testing that is the entire point of the design, and if it holds, Screen 1 should be cut before it ever ships broadly.

Decision criteria (pre-committed before launch)

If we see…	Then…
Mic tap rate flat	Discoverability failed, iterate on placement and copy. Do not conclude that the voice is unwanted.
Tap rate up, completion rate down	Dictation UX or ASR is failing, escalate to correction-first UX (Direction B) and raise WER with Applied AI.
Completion up, voice repeat flat	Voice works but is not worth repeating. The problem was never discoverability and **this bet is wrong.** Kill it and re-open the direction ranking.
Voice repeat up, total inputs flat	Cannibalisation, no business case. Do not ship on a vanity metric.
Voice repeat up and total inputs up	Ship. Expand to Gig & Field Workers (Phase 2).

Open Questions & Decisions Taken

Decisions taken log

#	Decision	Rationale	Stage
1	North Star = Total Voice Inputs / week	Decomposes cleanly into Active Voice Users × Inputs per Active User, so features can be mapped to breadth or depth.	Milestone 1
2	Segment = Vernacular Bharat	Largest reachable base with no strong AI rival; highest voice affinity; clear product-owned levers.	Milestone 2
3	Rescoped "Barrier Fixability" → "Product-Fixable"	The original criterion implied the product team could fix ASR quality. That is an Engineering / Applied-AI workstream.	Milestone 2
4	Replaced "Competitive Gap" → "Business Impact"	Competitive gap is a market observation, not a PM decision lever. Business impact is something a PM can legitimately argue.	Milestone 2
5	Direction A selected over B, C and D	Attacks the largest barrier (habit), carries no hard engineering dependency, ships and tests independently.	Milestone 3
6	Tap-to-fix chips folded into A	Cheapest slice of Direction B; consumes ASR output that already exists; directly addresses the top complaint.	Milestone 3
7	Primary metric = voice repeat, not activation	87% have already tried voice. An activation metric would measure a problem we do not have.	Milestone 3
8	No numeric targets before baseline	The survey measured attitudes, not behaviour. A target set today would be fabricated.	Milestone 3
9	4-arm experiment rather than bundled A/B	Attribution between discoverability and the habit loop is the decision-relevant question.	Milestone 3
10	Silent rollout recommended over launch comms	Nothing to announce until the experiment reads out.	Milestone 3

What has been descoped

Descoped	Why, and where it goes
Direction B: full correction-first UX	High impact and it addresses the top technical complaint, but it depends on ASR confidence-score access and ships more slowly alone. Minimal slice (tap-to-fix chips) folded into this scope; full scope is Phase 2.
Direction C: discreet / whisper mode	Addresses 48% who are socially constrained, but requires mic-sensitivity work owned by Applied AI, and the barrier is partly environmental rather than product-fixable.
Direction D: paywall friction	74% have never seen the limit message, so the addressable population is small. Handle as a separate low-effort quick win, not as this project.
Onboarding teaser (Screen 1) conditionally	If baseline confirms trial is genuinely high, Screen 1 solves a problem we do not have. Arm 2 of the experiment exists to find out. Flagged as a live cut candidate.
Segments: Students, Urban Professionals, Developers	Deprioritised in Milestone 2 on the scorecard due to environmental barrier, low voice affinity, and task mismatch respectively.

Descoped	Why, and where it goes
Segments: Gig & Field Workers; Older / Low-Literacy Adults	Phased to 2 and 3. Strong affinity but smaller base; and highest need but lowest reach.

Trade-offs made

1. **Choose the largest barrier over the loudest complaint.** Accuracy is cited more often as a discrete issue than habit appears in open-ended responses. We chose habit anyway, because it is the largest qualitative theme and decisively it is the one of the two we own. Reasonable people can disagree with this; it is the central bet of the document.
2. **Chose to ship an interruption to a segment that is 48% socially constrained.** The frequency governor mitigates it, but does not eliminate it. This is a real cost imposed on users like Kavita, and the dismiss-rate and session-abandonment guardrails exist specifically to catch it.
3. **Accepted a bounded ceiling.** If Hinglish WER stays high, no amount of nudging produces repeat use. This direction's upside is capped by a workstream we do not own. That is tracked as a dependency and stated openly rather than argued around.
4. **Choose attribution over speed.** A four-arm test costs sample size and calendar time versus a simple bundled A/B. Worth it, because "it worked" without "which part" leaves us unable to cut Screen 1.
5. **Choose Hinglish in Latin script for the wireframes, provisionally.** It matches how the segment writes today and sidesteps a font question, but it may be wrong for Kavita, who reads Devanagari more comfortably than Latin. Open question 2.

Open questions

#	Question	Why it matters	Owner	Needed by
1	Is the 87% trial rate real for the segment, or an artifact of LinkedIn sampling?	Decides whether Screen 1 is worth building at all, and whether the primary metric is right.	Data	W2
2	Hinglish in Latin script, or native Devanagari and regional scripts?	Affects every string in the feature and possibly the persona split.	Localisation	W2
3	Does Applied AI expose per-language WER reporting today?	Without it we cannot read our own ceiling or interpret a flat result.	Applied AI	W2
4	Are ASR confidence scores and n-best alternates exposed on the current API?	Blocking prerequisite for the correction chips (Screen 4).	Applied AI	W2
5	What is the acceptable nudge frequency ceiling before dismiss rate becomes a retention risk?	Determines the frequency governor's default.	Product + Data	W3
6	Is a streak mechanic appropriate for this segment, or does it read as gamification noise?	The reinforcement loop is the core of Arm 3. If the mechanic is wrong, so is the arm.	Design / Research	W3
7	Does mic-permission prompt timing in onboarding hurt completion?	62% of Indian users report AI-privacy concerns; asking too early may cost us the funnel.	Legal + Design	W3